

Project 3

Ahmed Tlili, Leon Petrinis, Mathilde Peruzzo

March 20, 2025

Relational Schema

- **Store**(s_id, s_address, phone_number, manager_id UNIQUE NOT NULL)
FOREIGN KEY(manager_id) REFERENCES Employee(employee_id)
- **Employee**(e_id, e_name, s_id)
FOREIGN KEY(s_id) REFERENCES Store(s_id)
- **Manufacturer**(m_id, m_name)
- **Product**(p_id, p_name NOT NULL, unit_price NOT NULL, description, discount_percentage, m_id NOT NULL)
FOREIGN KEY(m_id) REFERENCES Manufacturer(m_id)
- **Paint**(p_id, base, color)
FOREIGN KEY(p_id) REFERENCES Product(p_id)
- **Tool**(p_id, type)
- **Has_in_stock**(p_id, s_id, quantity NOT NULL)
FOREIGN KEY(p_id) REFERENCES Product(p_id)
FOREIGN KEY(s_id) REFERENCES Store(s_id)
- **Customer**(email, c_name, c_address NOT NULL)
PRIMARY KEY(email)
- **Purchase**(p_id, amount NOT NULL, p_date NOT NULL, p_time NOT NULL)
- **Contains_purchase**(p_id, product_id, quantity NOT NULL)
FOREIGN KEY(p_id) REFERENCES Purchase(p_id)
FOREIGN KEY(product_id) REFERENCES Product(p_id)
- **Instore**(p_id, e_id)
FOREIGN KEY(p_id) REFERENCES Purchase(p_id)
FOREIGN KEY(e_id) REFERENCES Employee(e_id)
- **Online**(p_id, rating, delivery_fee NOT NULL, email NOT NULL)
FOREIGN KEY(p_id) REFERENCES Purchase(p_id)
FOREIGN KEY(email) REFERENCES Customer(email)

Stored Procedure

Application Program

Indexing

Index 1

- (a) `db2 => CREATE INDEX clustered_purchase_idx ON Purchase(p_date) CLUSTER;
DB20000I The SQL command completed successfully.`

- (b) A clustered index on purchase date in the Purchase table is beneficial because purchases are frequently analyzed based on dates and date ranges. Thus, sorting the purchases by date allows for efficient range queries, making it faster to access data for accounting purposes. An example query that would benefit from this index is the following:

```
SELECT SUM(amount) AS total
FROM Purchase
WHERE p_date >= '01/01/2025' AND p_date <= '12/31/2025';
```

This above query computes the total revenue for the year 2025. With this clustered index, the database can quickly locate the first matching row, and perform a sequential scan to retrieve all rows within the specified date range, without needing to follow the pointers of other data entries (value + rid), as in a non-clustered index, which could often lead to more IO.

Index 2

- (a) `db2 => CREATE INDEX stock_idx ON Has_in_stock(s_id, quantity) CLUSTER;
DB20000I The SQL command completed successfully.`

- (b) This index is on the s_id and quantity attribute of the Has_in_stock table. It is useful for this application to efficiently identify products that are running low in stock in a specific store, which is crucial for inventory management. An example query that would benefit from this index is the following:

```
SELECT p_id
FROM Has_in_stock
WHERE s_id = 0 AND quantity < 5;
```

This query identifies all products of a particular store where the quantity of the product is very limited. The fact that it is a clustered index, again, allows for efficient range queries, making it faster to access data for inventory management purposes. It also makes sense to use a clustered index because the other attributes of the table are id's, which are certainly not needed in a sorted order.

Visualisation

Vis 1

Vis 2

Creativity